

# COMS4040A & COMS7045A: High Performance Computing & Scientific Data Management Introduction to MPI II

Hairong Wang

School of Computer Science & Applied Mathematics  
University of the Witwatersrand, Johannesburg

Semester one 2025

- 1 Point to Point Communication
  - Blocking vs. Non-blocking
  - MPI Message Passing Routine Arguments
  - Avoiding Deadlocks
  - Sending and Receiving Messages Simultaneously
  - Overlapping Communication with Computation
- 2 Collective Communications
  - MPI Collective Communication Illustrations
- 3 Examples
- 4 Summary

## 1 Point to Point Communication

- Blocking vs. Non-blocking
- MPI Message Passing Routine Arguments
- Avoiding Deadlocks
- Sending and Receiving Messages Simultaneously
- Overlapping Communication with Computation

## 2 Collective Communications

- MPI Collective Communication Illustrations

## 3 Examples

## 4 Summary

## 1 Point to Point Communication

- **Blocking vs. Non-blocking**
- MPI Message Passing Routine Arguments
- Avoiding Deadlocks
- Sending and Receiving Messages Simultaneously
- Overlapping Communication with Computation

## 2 Collective Communications

- MPI Collective Communication Illustrations


## 3 Examples

## 4 Summary

# Blocking vs. Non-blocking

- Most of the MPI point-to-point routines can be used in either blocking or non-blocking mode.
- Blocking:
  - A blocking send routine will only “return” after it is safe to modify the application buffer (your send data) for reuse.
  - A blocking send can be synchronous which means there is handshaking occurring with the receive task to confirm a safe send.
  - A blocking send can be asynchronous if a system buffer is used to hold the data for eventual delivery to the receive.
  - A blocking receive only “returns” after the data has arrived and is ready for use by the program.

# Blocking vs. Non-blocking cont.



**Figure:** Communication between two processes awakes both of them while transferring data from sender to receiver, possibly with a set of shorter sub-messages.

# Blocking vs. Non-blocking cont.

- Non-blocking:
  - Non-blocking send and receive routines behave similarly - they will return almost immediately. They do not wait for any communication events to complete.
  - Non-blocking operations simply “request” the MPI library to perform the operation when it is able.
  - It is unsafe to modify the application buffer (your variable space) until you know for a fact the requested non-blocking operation was actually performed by the library.
  - Non-blocking communications are primarily used to overlap computation with communication and exploit possible performance gains.

## 1 Point to Point Communication

- Blocking vs. Non-blocking
- **MPI Message Passing Routine Arguments**
- Avoiding Deadlocks
- Sending and Receiving Messages Simultaneously
- Overlapping Communication with Computation

## 2 Collective Communications

- MPI Collective Communication Illustrations

## 3 Examples

## 4 Summary



# MPI Message Passing Routine Arguments

MPI point-to-point communication routines generally have an argument list that takes one of the following formats:

---

Blocking sends	<code>MPI_Send(buffer, count, type, dest, tag, comm)</code>
Non-blocking sends	<code>MPI_Isend(buffer, count, type, dest, tag, comm, request)</code>
Blocking receive	<code>MPI_Recv(buffer, count, type, source, tag, comm, status)</code>
Non-blocking receive	<code>MPI_Irecv(buffer, count, type, source, tag, comm, request)</code>

---

## 1 Point to Point Communication

- Blocking vs. Non-blocking
- MPI Message Passing Routine Arguments
- **Avoiding Deadlocks**
- Sending and Receiving Messages Simultaneously
- Overlapping Communication with Computation

## 2 Collective Communications

- MPI Collective Communication Illustrations

## 3 Examples

## 4 Summary

# Avoiding Deadlocks

- The semantics of `MPI_Send` and `MPI_Recv` place some restrictions on how we can mix and match send and receive operations.
- Sources of Deadlocks:
  - Send a large message from process 0 to process 1
    - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
  - Mismatched send and receive – unsafe.
- What happens with
- Order the operations.

Process 0	Process 1
Send(1)	Send(0)
Recv(0)	Recv(1)

Process 0	Process 1
Send(1)	Recv(1)
Recv(0)	Send(0)

# Avoiding Deadlocks

## Example 1

Process 0 sends two messages with different tags to process 1, and process 1 receives them in reverse order.

```
1 int a[10], b[10], myrank;  
2 MPI_Status status;  
3 ...  
4 MPI_Comm_rank(MPI_COMM_WORLD, &myrank);  
5 if (myrank == 0) {  
6     MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);  
7     MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);  
8 } else if (myrank == 1) {  
9     MPI_Recv(b, 10, MPI_INT, 0, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
10    MPI_Recv(a, 10, MPI_INT, 0, 1, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
11 }  
12 ...
```

# Avoiding Deadlocks

## Example 2

Consider the following piece of code, in which process  $i$  sends a message to process  $i + 1$  (modulo the number of processes) and receives a message from process  $i - 1$  (modulo the number of processes).

```
int a[10], b[10], npes, myrank;
MPI_Status status;
...
MPI_Comm_size(MPI_COMM_WORLD, &npes);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
         MPI_COMM_WORLD);
MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
         MPI_COMM_WORLD, MPI_STATUS_IGNORE);
...
```

# Avoiding Deadlocks

## Example 2 cont.

We can break the circular wait to avoid deadlocks as follows:

```
int a[10], b[10], npes, myrank;
MPI_Status status;
...
MPI_Comm_size(MPI_COMM_WORLD, &npes);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank%2 == 1) {
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
             MPI_COMM_WORLD);
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
} else {
    MPI_Recv(b, 10, MPI_INT, (myrank-1+npes)%npes, 1,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Send(a, 10, MPI_INT, (myrank+1)%npes, 1,
             MPI_COMM_WORLD);
}
...
```

## 1 Point to Point Communication

- Blocking vs. Non-blocking
- MPI Message Passing Routine Arguments
- Avoiding Deadlocks
- **Sending and Receiving Messages Simultaneously**
- Overlapping Communication with Computation

## 2 Collective Communications

- MPI Collective Communication Illustrations

## 3 Examples

## 4 Summary

# Sending and Receiving Messages Simultaneously

To exchange messages, MPI provides the following function that both sends and receives a message:

```
int MPI_Sendrecv(void *sendbuf, int sendcount,  
                 MPI_Datatype senddatatype, int dest, int sendtag,  
                 void *recvbuf, int recvcount, MPI_Datatype recvdatatype,  
                 int source, int recvtag, MPI_Comm comm,  
                 MPI_Status *status)
```

The arguments include arguments to the send and receive functions.



# Using MPI\_Sendrecv in Example 2

Example 2 can be made “safe” by using MPI\_Sendrecv:

```
int a[10], b[10], npes, myrank;  
MPI_Status status;  
...  
MPI_Comm_size(MPI_COMM_WORLD, &npes);  
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);  
MPI_Sendrecv(a, 10, MPI_INT, (myrank+1)%npes, 1, b, 10,  
             MPI_INT, (myrank-1+npes)%npes, 1,  
             MPI_COMM_WORLD, &status);  
...
```

## 1 Point to Point Communication

- Blocking vs. Non-blocking
- MPI Message Passing Routine Arguments
- Avoiding Deadlocks
- Sending and Receiving Messages Simultaneously
- **Overlapping Communication with Computation**

## 2 Collective Communications

- MPI Collective Communication Illustrations

## 3 Examples

## 4 Summary

# Overlapping Communication with Computation

- In order to overlap communication with computation, MPI provides a pair of functions for performing non-blocking send and receive operations.

```
int MPI_Isend(void *buf, int count, MPI_Datatype datatype,  
             int dest, int tag, MPI_Comm comm,  
             MPI_Request *request)  
int MPI_Irecv(void *buf, int count, MPI_Datatype datatype,  
             int source, int tag, MPI_Comm comm,  
             MPI_Request *request)
```

- These operations return before the operations have been completed. Function `MPI_Test` tests whether or not the non-blocking send or receive operation identified by its `request` has finished.
- ```
int MPI_Test(MPI_Request *request, int *flag,  
            MPI_Status *status)
```

- **MPI\_Wait** blocks until a specified non-blocking send or receive operation has completed. For multiple non-blocking operations, the programmer can specify any, all or some completions.

```
int MPI_Wait(MPI_Request *request, MPI_Status *status)
```

- **MPI\_Request** handle is used to determine whether an operations has completed.
  - Non-blocking wait: **MPI\_Test**
  - Blocking wait: **MPI\_Wait**

- Anywhere you use **MPI\_Send** or **MPI\_Recv**, you can use the pair of **MPI\_Isend/MPI\_Wait** or **MPI\_Irecv/MPI\_Wait**.

- It is sometimes desirable to wait on multiple requests:

- **MPI\_Waitall**(int count,  
MPI\_Request array\_of\_requests[],  
MPI\_Status array\_of\_statuses[])

- The corresponding version of **MPI\_Test**

```
int MPI_Testall(int count,  
MPI_Request array_of_requests[], int *flag,  
MPI_Status array_of_statuses[])
```

flag: true if all the requests are completed, otherwise false

## Example 3

```
1 int main(int argc, char *argv[]){
2     int myid, numprocs, left, right, flag=0;
3     int buffer1[10], buffer2[10];
4     MPI_Request request; MPI_Status status;
5     MPI_Init(&argc,&argv);
6     MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
7     MPI_Comm_rank(MPI_COMM_WORLD, &myid);
8     /* initialize buffer2 */
9     .....
10    right = (myid + 1) % numprocs;
11    left = myid - 1;
12    if (left < 0)
13        left = numprocs - 1;
14    MPI_Irecv(buffer1, 10, MPI_INT, left, 123, MPI_COMM_WORLD,
15             &request);
16    MPI_Send(buffer2, 10, MPI_INT, right, 123, MPI_COMM_WORLD);
17    MPI_Test(&request, &flag, &status);
18    while (!flag){
19        /* Do some work ... */
20        MPI_Test(&request, &flag, &status);
21    }
22    MPI_Finalize();
23 }
```

## Example 4


```
1 int main(int argc, char *argv){
2     int numtasks, rank, next, prev, buf[2], tag1=1, tag2=2;
3     MPI_Request reqs[4]; MPI_Status stats[4];
4     MPI_Init(&argc,&argv);
5     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
6     MPI_Comm_rank(MPI_COMM_WORLD, &rank);

7
8     prev = rank-1; next = rank+1;
9     if (rank == 0) prev = numtasks - 1;
10    if (rank == (numtasks - 1)) next = 0;
11    MPI_Irecv(&buf[0],1,MPI_INT,prev,tag1,MPI_COMM_WORLD,
12             &reqs[0]);
13    MPI_Irecv(&buf[1],1,MPI_INT,next,tag2,MPI_COMM_WORLD,
14             &reqs[1]);

15
16    MPI_Isend(&rank,1,MPI_INT,prev,tag2,MPI_COMM_WORLD,
17             &reqs[2]);
18    MPI_Isend(&rank,1,MPI_INT,next,tag1,MPI_COMM_WORLD,
19             &reqs[3]);
20    MPI_Waitall(4, reqs, stats);
21    MPI_Finalize();
22 }
```

## Example 5 (The Trapezoidal Rule)

- We can use **the trapezoidal rule** to approximate the area between the graph of a function,  $y = f(x)$ , two vertical lines, and the  $x$ -axis.



**Figure:** The trapezoidal rule: (a) area to be estimated, (b) estimate area using trapezoids

### Example 5 cont.

- If the endpoints of the subinterval are  $x_i$  and  $x_{i+1}$ , then the length of the subinterval is  $h = x_{i+1} - x_i$ . Also, if the lengths of the two vertical segments are  $f(x_i)$  and  $f(x_{i+1})$ , then the area of the trapezoid is

$$\text{Area of one trapezoid} = \frac{h}{2}(f(x_i) + f(x_{i+1})).$$

- Since we chose the  $N$  subintervals, we also know that the bounds of the region are  $x = a$  and  $x = b$  then

$$h = \frac{b - a}{N}$$



## Example 5 cont.

- The pseudo code for a serial program:

```
h = (b-a) / N;  
approx = (f(a) + f(b)) / 2.0;  
for(i=1; i<=n-1; i++){  
    x_i = a + i * h;  
    approx += f(x_i);  
}  
approx = h * approx;
```

Recall we can design a parallel program using four basic steps:

- 1 Partition the problem solution into tasks.
- 2 Identify the communication between the tasks.
- 3 Aggregate the tasks into composite tasks.
- 4 Map the composite tasks to cores.

## Example 5: Parallel Algorithm for the Trapezoidal Rule

Assuming `comm_sz` evenly divides  $n$ , the pseudo-code for the parallel program looks like the following:

```
1  Get a, b, n;
2  h = (b - a)/n;
3  local_n = n/comm_sz;
4  local_a = a + my_rank * local_n * h;
5  local_b = local_a + local_n * h;
6  local_integral = Trap(local_a, local_b, local_n, h);
7  if (my_rank != 0)
8      Send local integral to process 0;
9  else { /* my_rank == 0 */
10     total_integral = local_integral;
11     for (proc = 1; proc < comm_sz; proc++) {
12         Receive local_integral from proc;
13         total_integral += local_integral;
14     }
15 }
16 if (my_rank == 0)
17     print result;
```

# Dealing with I/O

- In most cases, all the processes in `MPI_COMM_WORLD` have access to `stdout` and `stderr`.
- The order in which the processes' output appears is indeterministic.
- For the input, i.e., `stdin`, usually, only process 0 has access to.
- If an MPI program uses `scanf` function, then process 0 reads in the data, and sends it to the other processes.

# Outline

- 1 Point to Point Communication
  - Blocking vs. Non-blocking
  - MPI Message Passing Routine Arguments
  - Avoiding Deadlocks
  - Sending and Receiving Messages Simultaneously
  - Overlapping Communication with Computation
- 2 **Collective Communications**
  - MPI Collective Communication Illustrations
- 3 Examples
- 4 Summary

# Collective Communications

- Communication is coordinated among a group of processes, as specified by a communicator.
- All collective operations are blocking and no message tags are used.
- All processes in the communicator must call the collective operation.
- Three classes of collective operations
  - Data movement
  - Collective computation
  - Synchronization

# Outline

- 1 Point to Point Communication
  - Blocking vs. Non-blocking
  - MPI Message Passing Routine Arguments
  - Avoiding Deadlocks
  - Sending and Receiving Messages Simultaneously
  - Overlapping Communication with Computation
- 2 **Collective Communications**
  - MPI Collective Communication Illustrations
- 3 Examples
- 4 Summary

# MPI\_Bcast


- A collective communication in which data belonging to a single process is sent to all of the processes in the communicator is called a broadcast — `MPI_Bcast`.

## MPI\_Bcast

Broadcasts a message to all other processes of that group

```
count = 1;
source = 1;
MPI_Bcast(msg, count, MPI_INT, source, MPI_COMM_WORLD);
```

broadcast originates in task 1



- The process with rank `source` sends the contents of the memory referenced by `msg` to all the processes in the communicator `MPI_COMM_WORLD`.

## Example 5 cont.

- 1 In the example `mpi_trapezoid_1.c`, we are using

```
1 if (my_rank == 0) {
2     for (dest = 1; dest < comm_sz; dest++) {
3         MPI_Send(a_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
4         MPI_Send(b_p, 1, MPI_DOUBLE, dest, 0, MPI_COMM_WORLD);
5         MPI_Send(n_p, 1, MPI_INT, dest, 0, MPI_COMM_WORLD);
6     } else { /* my rank != 0 */
7         MPI_Recv(a_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD,
8                 MPI_STATUS_IGNORE);
9         MPI_Recv(b_p, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD,
10                MPI_STATUS_IGNORE);
11        MPI_Recv(n_p, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
12               MPI_STATUS_IGNORE);
13    }
14 }
```

- 2 Instead of using point-to-point communications, you can use collective communications here. **Write another function to implement this part using** — `MPI_Bcast()`.




# MPI\_Reduce

- `MPI_Reduce` combines data from all processes in the communicator and returns it to one process.
- In many numerical algorithms, `Send/Receive` can be replaced by `Bcast/Reduce`, improving both simplicity and efficiency.

## MPI\_Reduce

Perform and associate reduction operation across all tasks in the group and place the result in one task

```
count = 1;  
dest = 1;           result will be placed in task 1  
MPI_Reduce(sendbuf, rcvbuf, count, MPI_INT, MPI_SUM,  
            dest, MPI_COMM_WORLD);
```



# MPI\_Reduce cont.

- When the `count` is greater 1, `MPI_Reduce` operate on arrays instead of scalars.

```
1 double local_x[N], sum[N];  
2 ...  
3 MPI_Reduce(local_x, sum, N, MPI_DOUBLE, MPI_SUM, 0, \  
4           MPI_COMM_WORLD);
```

# Example 5 cont.

- ❶ In the example `mpi_trapezoid_1.c`, we are using

```
1  /* Add up the integrals calculated by each process */
2  if (my_rank != 0) {
3      MPI_Send(&local_int, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
4  } else {
5      total_int = local_int;
6      for (source = 1; source < comm_sz; source++) {
7          MPI_Recv(&local_int, 1, MPI_DOUBLE, source, 0,
8                  MPI_COMM_WORLD, MPI_STATUS_IGNORE);
9          total_int += local_int;
10     }
```

- ❷ Instead of using point-to-point communications, you can also use collective communications here. Rewrite this part using appropriate collective communication.

# MPI\_Reduce

- Suppose that each process calls `MPI_Reduce` with operator `MPI_SUM`, and **destination process 0**. What happens with the following multiple calls of `MPI_Reduce`? What are the values for `b` and `d`?

| Time | Process 0                                     | Process 1                                     | Process 2                                     |
|------|-----------------------------------------------|-----------------------------------------------|-----------------------------------------------|
| 0    | <code>a=1; c = 2;</code>                      | <code>a=1; c = 2;</code>                      | <code>a=1; c = 2;</code>                      |
| 1    | <code>MPI_Reduce (&amp;a, &amp;b, ...)</code> | <code>MPI_Reduce (&amp;c, &amp;d, ...)</code> | <code>MPI_Reduce (&amp;a, &amp;b, ...)</code> |
| 2    | <code>MPI_Reduce (&amp;c, &amp;d, ...)</code> | <code>MPI_Reduce (&amp;a, &amp;b, ...)</code> | <code>MPI_Reduce (&amp;c, &amp;d, ...)</code> |

- The order of the calls will determine the matching.
- What will happen with the following code?

```
MPI_Reduce (&x, &x, 1, MPI_DOUBLE, MPI_SUM, 0, comm);
```

# MPI\_Allreduce

- If the result of the reduction operation is needed by all processes, MPI provides:


```
int MPI_Allreduce(void *sendbuf, void *recvbuf,  
    int count, MPI_Datatype datatype, MPI_Op op,  
    MPI_Comm comm)
```

- This is equivalent to an MPI\_Reduce followed by an MPI\_Bcast.


## MPI\_Allreduce

Perform and associate reduction operation across all tasks in the group and place the result in all tasks


```
count = 1;  
MPI_Allreduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,  
    MPI_COMM_WORLD);
```



# MPI\_Allreduce cont.



(a)



(b)

**Figure:** (a) A global sum followed by a broadcasting; (2) A butterfly structured global sum.

- The scatter operation is to distribute **distinct messages** from a single source task to each task in the group.


```
int MPI_Scatter(void *sendbuf, int sendcount,  
               MPI_Datatype senddatatype, void *recvbuf,  
               int recvcount, MPI_Datatype recvdatatype,  
               int source, MPI_Comm comm)
```

## MPI\_Scatter

Sends data from one task to all other tasks in a group

```
sendcnt = 1;  
recvcnt = 1;  
src = 1;  
MPI_Scatter(sendbuf, sendcnt, MPI_INT,  
            recvbuf, recvcnt, MPI_INT,  
            src, MPI_COMM_WORLD);
```

task 1 contains the message to be scattered





- The gather operation is performed in MPI using `MPI_Gather`.
  - Gathers **distinct messages** from each task in the group to a single destination task.
  - Reverse operation of `MPI_Scatter`.


```
int MPI_Gather(void *sendbuf, int sendcount,  
              MPI_Datatype senddatatype, void *recvbuf,  
              int recvcnt, MPI_Datatype recvdatatype,  
              int target, MPI_Comm comm)
```

# MPI\_Gather cont.

## MPI\_Gather

Gathers together values from a group of processes

```
sendcnt = 1;  
recvcnt = 1;  
src = 1;          messages will be gathered in task 1  
MPI_Gather(sendbuf, sendcnt, MPI_INT,  
           rcvbuf, recvcnt, MPI_INT,  
           src, MPI_COMM_WORLD);
```



- MPI also provides the `MPI_Allgather` function in which the data are gathered at all the processes.


```
int MPI_Allgather(void *sendbuf, int sendcount,  
                 MPI_Datatype senddatatype, void *recvbuf,  
                 int recvcount, MPI_Datatype recvdatatype,  
                 MPI_Comm comm)
```

# MPI\_Allgather cont.

## MPI\_Allgather

Gathers together values from a group of processes and distributes to all

```
sendcnt = 1;  
recvcnt = 1;  
MPI_Allgather(sendbuf, sendcnt, MPI_INT,  
              recvbuf, recvcnt, MPI_INT,  
              MPI_COMM_WORLD);
```



- The all-to-all communication operation is performed by:

```
int MPI_Alltoall(void *sendbuf, int sendcount,  
                MPI_Datatype senddatatype, void *recvbuf,  
                int recvcount, MPI_Datatype recvdatatype,  
                MPI_Comm comm)
```

- Each task in a group performs a scatter operation, sending a distinct message to all the tasks in the group in order by index.

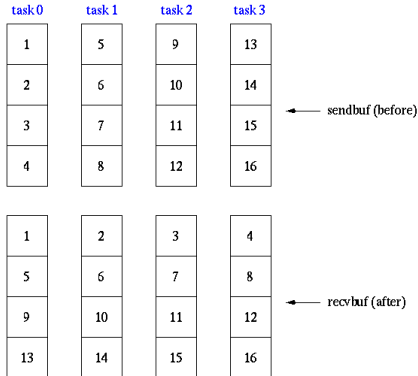
# MPI\_Alltoall cont.

## MPI\_Alltoall

Sends data from all to all processes. Each process performs a scatter operation.

```
sendcnt = 1;  
recvcnt = 1;
```

```
MPI_Alltoall(sendbuf, sendcnt, MPI_INT,  
             recvbuf, recvcnt, MPI_INT,  
             MPI_COMM_WORLD);
```



## Example 6 (Matrix vector multiplication)

If  $A = (a_{ij})$  is an  $m \times n$  matrix and  $\mathbf{x}$  is a vector with  $n$  components, then  $\mathbf{y} = A\mathbf{x}$  is a vector with  $m$  components. Furthermore,

$$y_i = a_{i0}x_0 + a_{i1}x_1 + a_{i2}x_2 + \dots + a_{i,n-1}x_{n-1}.$$

A serial code can be as simple as

```
1  for (i = 0; i < m; i++) {  
2      y[i] = 0.0;  
3      for (j = 0; j < n; j++)  
4          y[i] += A[i*n+j]*x[j];  
5  }
```

## Example 6 cont.

Process 0 reads in the matrix and distributes row blocks to all the processes in communicator `comm`.

```
1 if (my_rank == 0) {
2   A = malloc(m*n*sizeof(double));
3   if (A == NULL) local_ok = 0;
4   Check_for_error(local_ok, "Random_matrix",
5     "Can't allocate temporary matrix", comm);
6   srand(2018);
7   for (i = 0; i < m; i++)
8     for (j = 0; j < n; j++)
9       A[i*n+j] = (double)rand() / RAND_MAX;
10  MPI_Scatter(A, local_m*n, MPI_DOUBLE,
11    local_A, local_m*n, MPI_DOUBLE, 0, comm);
12  free(A);
13 } else {
14   Check_for_error(local_ok, "Random_matrix",
15     "Can't allocate temporary matrix", comm);
16   MPI_Scatter(A, local_m*n, MPI_DOUBLE,
17     local_A, local_m*n, MPI_DOUBLE, 0, comm);
18 }
```



## Example 6 cont.

Each process gathers the entire vector, then proceeds to compute its share of sub-matrix and vector multiplication.

```
1 MPI_Allgather(local_x, local_n, MPI_DOUBLE,  
2   x, local_n, MPI_DOUBLE, comm);  
  
4 for (local_i = 0; local_i < local_m; local_i++) {  
5     local_y[local_i] = 0.0;  
6     for (j = 0; j < n; j++)  
7         local_y[local_i] += local_A[local_i*n+j]*x[j];  
8 }
```

# Exercise: Matrix Transpose

Implement matrix transpose using MPI scatter and gather operations.

# MPI\_Scan


- To compute prefix-sums, MPI provides:

```
int MPI_Scan(void *sendbuf, void *recvbuf, int count,  
            MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

## MPI\_Scan

Computes the scan (partial reductions) of data on a collection of processes

```
count = 1;  
MPI_Scan(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,  
        MPI_COMM_WORLD);
```



- Using this core set of collective operations, a number of programs can be greatly simplified.

Scatters a buffer in parts to all processes in a communicator, which allows different amounts of data to be sent to different processes.

```
int MPI_Scatterv(void *sendbuf, int *sendcounts, int *displs,  
MPI_Datatype sendtype, void *recvbuf, int recvcount,  
MPI_Datatype recvtype, int source, MPI_Comm comm)
```

- **sendbuf**: address of send buffer (significant only at root)
- **sendcounts**: integer array (of length group size) specifying the number of elements to send to each processor
- **displs**: integer array (of length group size). Entry *i* specifies the displacement (relative to **sendbuf** from which to take the outgoing data to process *i*)
- **sendtype**: data type of send buffer elements
- **recvcount**: number of elements in receive buffer (integer)
- **recvtype**: data type of receive buffer elements
- **root**: rank of sending process (integer)

# MPI\_Scatterv cont.

## Example 7

Given an  $N \times N$  matrix,  $A$ , of integers, write an MPI program that distributes the first  $M$  rows of the upper triangle of  $A$  to  $M$  processes by rows, where each process gets one row of the upper triangle of  $A$  (when  $M = N$ , it means each process gets one row of the upper triangle of  $A$ ).

For this example, we can use `MPI_Scatterv`.

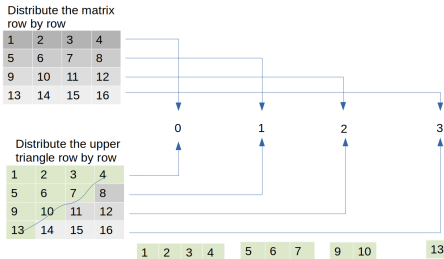


Figure: `MPI_Scatter` (top) and `MPI_Scatterv` (bottom) example

# MPI\_Scatterv cont.

For Example 7, assuming the matrix is only  $4 \times 4$ , and we are running the MPI code using 4 processes, then some of the arguments of calling `MPI_Scatterv`:

- `sendcounts[4] = {4, 3, 2, 1};`
- `displs[4] = {0, 4, 8, 12}` which is with reference to `sendbuf`; these values can be expressed as  $N * \text{rank}$ , where `rank` is the rank of a process.
- note also that `recvcount` in `MPI_Scatterv` is a scalar; for process 0, `recvcount = 4 (=4-0)`; for process 1, `recvcount = 3 (=4-1)`; for process 2, `recvcount = 2 (=4-2)`; and for process 3, `recvcount = 1 (=4-3)`; so this value can be obtained as  $N - \text{rank}$  where  $N$  is the number of rows in the matrix, and `rank` is the rank of a process.

`scatterv_1.c` gives an example code for Example 7.

Sends data from all to all processes; each process may send a different amount of data and provide displacements for the input and output data.

```
MPI_Alltoallv(void *sendbuf, int *sendcounts, int *sdispls,  
MPI_Datatype sendtype, void *recvbuf, int *recvcounts,  
int *rdispls, MPI_Datatype recvtype, MPI_Comm comm)
```

- **sendbuf**: starting address of send buffer
- **sendcounts**: integer array equal to the group size specifying the number of elements to send to each processor
- **sdispls**: integer array (of length group size). Entry *j* specifies the displacement (relative to **sendbuf** from which to take the outgoing data destined for process *j*)
- **sendtype**: data type of send buffer elements
- **recvcounts**: integer array equal to the group size specifying the maximum number of elements that can be received from each processor
- **rdispls**: integer array (of length group size). Entry *i* specifies the displacement (relative to **recvbuf** at which to place the incoming data from process *i*)
- **recvtype**: data type of receive buffer elements

## Example 8

Given the `MPI_Alltoallv` argument settings shown in the figure (the number of processes is 3), what is the content of `recvbuf` for each process?

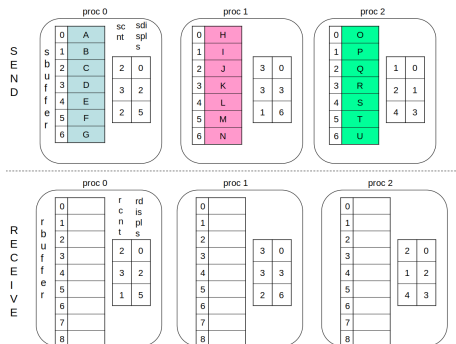


Figure: `MPI_Alltoallv` example



# MPI\_Alltoallv cont.

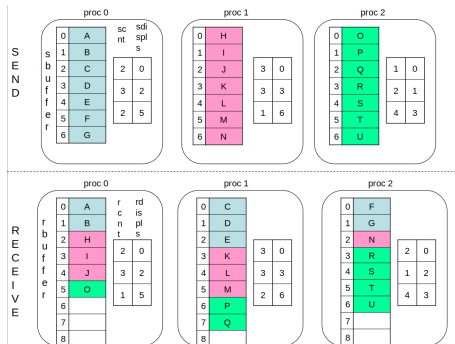


Figure: MPI\_Alltoallv example

# MPI\_Gatherv

The following function allows a different number of data elements to be sent by each process by replacing `recvcount` in `MPI_Gather` with an array `recvcounts`

```
int MPI_Gatherv(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
void *recvbuf, int *recvcounts, int *displs,  
MPI_Datatype recvtpe, int target, MPI_Comm comm)
```

- `sendbuf`: pointer, starting address of send buffer (or the data to be sent)
- `sendcount`: the number of elements in the send buffer
- `sendtype`: datatype of send buffer elements
- `recvbuf`: pointer, starting address of receive buffer (significant only at root)
- `recvcounts`: integer array (of length group size) containing the number of elements to be received from each process (significant only at root)
- `displs`: integer array (of length group size). Entry `i` specifies the displacement relative to `recvbuf` at which to place the incoming data from process `i` (significant only at root)
- `recvtpe`: the datatype of data to be received (significant only at root)
- `target`: rank of receiving process (integer)

# MPI\_Allgatherv

Gather data from all processes and deliver the combined data to all processes

```
int MPI_Allgatherv(void *sendbuf, int sendcount, MPI_Datatype sendtype,  
void *recvbuf, int *recvcounts, int *displs,  
MPI_Datatype recvtype, MPI_Comm comm)
```

- **sendbuf**: pointer, starting address of send buffer (or the data to be sent)
- **sendcount**: the number of elements in the send buffer
- **sendtype**: datatype of send buffer elements
- **recvbuf**: pointer, starting address of receive buffer (significant only at root)
- **recvcounts**: integer array (of length group size) containing the number of elements to be received from each process (significant only at root)
- **displs**: integer array (of length group size). Entry *i* specifies the displacement relative to **recvbuf** at which to place the incoming data from process *i* (significant only at root)
- **recvtype**: the datatype of data to be received (significant only at root)


# Outline

- 1 Point to Point Communication
  - Blocking vs. Non-blocking
  - MPI Message Passing Routine Arguments
  - Avoiding Deadlocks
  - Sending and Receiving Messages Simultaneously
  - Overlapping Communication with Computation
- 2 Collective Communications
  - MPI Collective Communication Illustrations
- 3 Examples
- 4 Summary

# Parallel quicksort

- one process broadcast initial pivot to all processes;
- each process in the upper half swaps with a partner in the lower half
- recurse on each half
- swap among partners in each half
- each process uses quicksort on local elements

# Parallel quicksort cont.




# Hyperquicksort

Limitation of parallel quicksort: poor balancing of list sizes.

Hyperquicksort: sort elements before broadcasting pivot.

- sort elements in each process
- select median as pivot element and broadcast it
- each process in the upper half swaps with a partner in the lower half
- recurse on each half

# Hyperquicksort cont.





## Example 9 (Task 0 pings task 1 and awaits return ping)

```
1 #include "mpi.h"
2 #include <stdio.h>
3 int main(int argc, char *argv[]) {
4     int numtasks, rank, dest, source, rc, count, tag=1;
5     char inmsg, outmsg='x';
6     MPI_Status Stat;
7     MPI_Init(&argc,&argv);
8     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
9     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
10
11     if (rank == 0){
12         dest = 1;
13         source = 1;
14         rc = MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
15         rc = MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &
16             Stat);
17     }
18     else if (rank == 1) {
19         dest = 0;
20         source = 0;
21         rc = MPI_Recv(&inmsg, 1, MPI_CHAR, source, tag, MPI_COMM_WORLD, &
22             Stat);
23         rc = MPI_Send(&outmsg, 1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
24     }
25     rc = MPI_Get_count(&Stat, MPI_CHAR, &count);
26     printf("Task %d: Received %d char(s) from task %d with tag %d \n",
27         rank, count, Stat.MPI_SOURCE, Stat.MPI_TAG);
28     MPI_Finalize();
29 }
```

## Example 10 (Nearest neighbor exchange in a ring topology)

```
1 #include "mpi.h"
2 #include <stdio.h>
3 int main(int argc, char *argv[]){
4     int numtasks, rank, next, prev, buf[2], tag1=1, tag2=2;
5     MPI_Request reqs[4];
6     MPI_Status stats[4];

7
8     MPI_Init(&argc,&argv);
9     MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
10    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

11
12    prev = rank-1;
13    next = rank+1;
14    if (rank == 0) prev = numtasks - 1;
15    if (rank == (numtasks - 1)) next = 0;
16    MPI_Irecv(&buf[0], 1, MPI_INT, prev, tag1, MPI_COMM_WORLD, &reqs[0]);
17    MPI_Irecv(&buf[1], 1, MPI_INT, next, tag2, MPI_COMM_WORLD, &reqs[1]);

18
19    MPI_Isend(&rank, 1, MPI_INT, prev, tag2, MPI_COMM_WORLD, &reqs[2]);
20    MPI_Isend(&rank, 1, MPI_INT, next, tag1, MPI_COMM_WORLD, &reqs[3]);

21
22    MPI_Waitall(4, reqs, stats);

23
24    MPI_Finalize();
25 }
```

## Example 11 (Perform a scatter operation on the rows of an array)

```
1 #include "mpi.h"
2 #include <stdio.h>
3 #define SIZE 4

5 int main(int argc, char *argv[]) {
6     int numtasks, rank, sendcount, recvcount, source;
7     float sendbuf[SIZE][SIZE] = {
8         {1.0, 2.0, 3.0, 4.0},
9         {5.0, 6.0, 7.0, 8.0},
10        {9.0, 10.0, 11.0, 12.0},
11        {13.0, 14.0, 15.0, 16.0} };
12    float recvbuf[SIZE];
13    MPI_Init(&argc,&argv);
14    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
15    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
16    if (numtasks == SIZE) {
17        source = 1;
18        sendcount = SIZE;
19        recvcount = SIZE;
20        MPI_Scatter(sendbuf,sendcount,MPI_FLOAT,recvbuf,recvcount,
21                  MPI_FLOAT,source,MPI_COMM_WORLD);

23        printf("rank= %d Results: %f %f %f %f\n",rank,recvbuf[0],
24              recvbuf[1],recvbuf[2],recvbuf[3]);
25    }
26    else
27        printf("Must specify %d processors. Terminating.\n",SIZE);

29    MPI_Finalize();
30 }
```

## Example 12 (The Odd-Even Transposition Sort)

- Sorts  $n$  elements in  $n$  phases ( $n$  is even), each of which requires  $n/2$  compare-exchange operations.
- The algorithm alternates between two phases — odd and even phases.
- Let  $\langle a_0, a_1, \dots, a_{n-1} \rangle$  be the sequence to be sorted.
  - During the odd phase, elements with odd indices are compared with their right neighbours, and if they are out of sequence they are exchanged; thus, the pairs  $(a_1, a_2), (a_3, a_4), \dots, (a_{n-3}, a_{n-2})$  are compare exchanged.
  - During the even phase, elements with even indices are compared with their right neighbours, and if they are out of sequence they are exchanged;  $(a_0, a_1), (a_2, a_3), \dots, (a_{n-2}, a_{n-1})$ .
- After  $n$  phases of odd-even exchanges, the sequence is sorted. Each phase requires  $n/2$  compare-exchange operations (sequential complexity  $O(n^2)$ ).

## Example 12 cont. – The serial algorithm

```
1 for i = 0 to n-1 do
2   if i is even then
3     for j = 0 to n/2 - 1 do
4       compare-exchange(a(2j), a(2j+1));
5   if i is odd then
6     for j = 0 to n/2 - 1 do
7       compare-exchange(a(2j+1), a(2j+2));
```

## Example 12 cont. – The parallel algorithm

```
1 void oddevensort(int n)
2   id = process's label;
3   for i =0 to n-1 do
4     if i is odd then
5       if id is odd then
6         compare-exchange_min(id, id + 1); //increasing comparator
7       else
8         compare-exchange_max(id, id - 1); //decreasing comparator
9     if i is even then
10      if id is even then
11        compare-exchange_min(id, id + 1);
12      else
13        compare-exchange_max(id, id - 1);
```

# Outline

- 1 Point to Point Communication
  - Blocking vs. Non-blocking
  - MPI Message Passing Routine Arguments
  - Avoiding Deadlocks
  - Sending and Receiving Messages Simultaneously
  - Overlapping Communication with Computation
- 2 Collective Communications
  - MPI Collective Communication Illustrations
- 3 Examples
- 4 Summary

# Summary





- Point-to-point communication
  - Blocking vs non-blocking
  - Safety in MPI programs
- Collective communication
  - Collective communications involve all the processes in a communicator.
  - All the processes in the communicator must call the same collective function.
  - Collective communications do not use **tags**, the message is matched on the order in which they are called within the communicator.
  - The meanings of *local variable* and *global variable* in MPI
  - Some important MPI collective communications we learned:  
MPI\_Reduce, MPI\_Allreduce, MPI\_Bcast, MPI\_Gather,  
MPI\_Scatter, MPI\_Allgather, MPI\_Alltoall, MPI\_Scan  
etc.



The resources used include

- *Introduction to Parallel Computing*
- *MPI Forum*
- *Using MPI: Portable Parallel Programming with the Message Passing Interface*
- *Parallel Programming in C with MPI and OpenMP*
- <https://hpc-tutorials.llnl.gov/mpi/>

# References cont. I

-  Grama, Ananth et al. *Introduction to Parallel Computing*. Addison Wesley, 2003.
-  Gropp, William, Ewing Lusk, and Anthony Skjellum. *Using MPI: Portable Parallel Programming with the Message Passing Interface*. Second. The MIT Press, Cambridge, Massachusetts, London, England., 1999.
-  MPI. *MPI Forum*. <https://www.mpi-forum.org/docs/>. Accessed 2025-4-30. 2022.
-  Quinn, Michael J. *Parallel Programming in C with MPI and OpenMP*. McGraw-Hill Education Group, 2003.